

EMOD MPI Test Report

Zhaowei Du

(April 2024)

1 Table of Contents

EMOD MPI Test Report	1
1 Overview.....	3
2 Scenarios and results.....	4
2.1 Test Cases	4
2.2 Test Results.....	5
2.2.1 Cases: mpirun -n 1 & #SBATCH --ntasks=1	5
2.2.2 Cases: mpirun -n 2 & #SBATCH --ntasks=2	6
2.2.3 Cases: mpirun -n 3 & #SBATCH --ntasks=3	7
2.2.4 Cases: mpirun -n 4 & #SBATCH --ntasks=4	8
2.3 Special test.....	9
3 Summary.....	10

1 Overview

The following is a typical command to run simulation with EMOD on SLURM:

```
singularity exec <user_sif_file> Assets/Eradication --config config.json --dll-path ./Assets
--input-path ./Assets\;
```

NYU user would like to efficiently use system cores to improve performance, SLURM has several related parameters, such as the following, user is not sure how to choose from:

```
#SBATCH --nodes=3
#SBATCH --ntasks=5
#SBATCH --cpus-per-task=2
#SBATCH --ntasks-per-core=4
```

Instead of using above SLURM parameters, user did a test by inserting 'mpirun -n #' to the above command as follows, and wonders if it is a good solution:


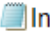
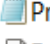
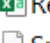
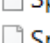
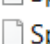
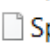
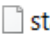
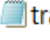

```
singularity exec <user_sif_file> mpirun -n 4 Assets/Eradication --config config.json --dll-
path ./Assets --input-path ./Assets\;
```

This document is a report based on my test on all kinds of scenarios. All the tests have been done on NU SLURM system.

We run the **same** experiment/simulations in different scenarios.

Note: our running test's demographic file has **node_count = 3**.

The following is the normal case result without 'mpirun -n #' being inserted:

			4/17/2024 1:18:47 PM
 InsetChart.json	7,862 KB		4/17/2024 1:35:46 PM
 PropertyReport.json	1,027 KB		4/17/2024 1:35:46 PM
 ReportEventRecorder.csv	9,079 KB		4/17/2024 1:35:43 PM
 SpatialReport_Adult_Vectors.bin	257 KB		4/17/2024 1:35:46 PM
 SpatialReport_Daily_Bites_Per_Human.bin	257 KB		4/17/2024 1:35:46 PM
 SpatialReport_Population.bin	257 KB		4/17/2024 1:35:46 PM
 SpatialReport_Prevalence.bin	257 KB		4/17/2024 1:35:46 PM
 state-21900.dtk	58,360 KB		4/17/2024 1:35:45 PM
 transitions.json	1 KB		4/17/2024 1:18:47 PM

We will use this to compare the results generated from each scenarios test.

2 Scenarios and results

On NU system, tested the following scenarios.

2.1 Test Cases

- (1) singularity exec <user_sif_file> mpirun -n # Assets/Eradication --config config.json --dll-path ./Assets --input-path ./Assets\;.

Case: mpirun -n 1
 Case: mpirun -n 2
 Case: mpirun -n 3
 Case: mpirun -n 4

Note: tested mpiexec as well.

- (2) mpirun -n # singularity exec <user_sif_file> Assets/Eradication --config config.json --dll-path ./Assets --input-path ./Assets\;.

Case: mpirun -n 1
 Case: mpirun -n 2
 Case: mpirun -n 3
 Case: mpirun -n 4

Note: tested mpiexec as well.

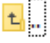
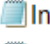
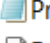
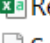
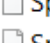
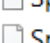
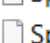
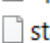
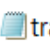

- (3) COMPS approach
 #!/usr/bin/bash
 #SBATCH --ntasks=<Number of instances to run>
 Srun -mpi=pmi2 singularity exec <user_sif_file> Assets/Eradication --config config.json --dll-path ./Assets --input-path ./Assets\;.

Case: #SBATCH --ntasks=1
 Case: #SBATCH --ntasks=2
 Case: #SBATCH --ntasks=3
 Case: #SBATCH --ntasks=4

Remark: the tests show that all above approaches work and behavior the same way. Refer to the next Section for details.

2.2 Test Results

Expected results:


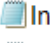
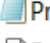
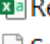
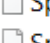
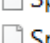
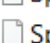
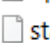
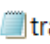

			4/17/2024 1:18:47 PM
	InsetChart.json	7,862 KB	4/17/2024 1:35:46 PM
	PropertyReport.json	1,027 KB	4/17/2024 1:35:46 PM
	ReportEventRecorder.csv	9,079 KB	4/17/2024 1:35:43 PM
	SpatialReport_Adult_Vectors.bin	257 KB	4/17/2024 1:35:46 PM
	SpatialReport_Daily_Bites_Per_Human.bin	257 KB	4/17/2024 1:35:46 PM
	SpatialReport_Population.bin	257 KB	4/17/2024 1:35:46 PM
	SpatialReport_Prevalence.bin	257 KB	4/17/2024 1:35:46 PM
	state-21900.dtk	58,360 KB	4/17/2024 1:35:45 PM
	transitions.json	1 KB	4/17/2024 1:18:47 PM

2.2.1 Cases: mpirun -n 1 & #SBATCH --ntasks=1

Results

- Run succeeded.
- Output as expected.

Example output

			4/17/2024 12:56:15 PM
	InsetChart.json	7,862 KB	4/17/2024 1:12:07 PM
	PropertyReport.json	1,027 KB	4/17/2024 1:12:07 PM
	ReportEventRecorder.csv	9,079 KB	4/17/2024 1:12:03 PM
	SpatialReport_Adult_Vectors.bin	257 KB	4/17/2024 1:12:07 PM
	SpatialReport_Daily_Bites_Per_Human.bin	257 KB	4/17/2024 1:12:07 PM
	SpatialReport_Population.bin	257 KB	4/17/2024 1:12:07 PM
	SpatialReport_Prevalence.bin	257 KB	4/17/2024 1:12:07 PM
	state-21900.dtk	58,360 KB	4/17/2024 1:12:06 PM
	transitions.json	1 KB	4/17/2024 12:56:15 PM

stdout.txt

...

```
00:15:48 [0] [I] [Simulation] Update(): Time: 21900.0 Rank: 0 StatPop: 3310 Infected: 397
00:15:48 [0] [I] [SerializedPopulation] Writing state to 'output/state-21900.dtk'
00:15:51 [0] [I] [Simulation] Finalizing 'InsetChart.json' reporter.
00:15:52 [0] [I] [Simulation] Finalized 'InsetChart.json' reporter.
00:15:52 [0] [I] [Simulation] Finalizing 'PropertyReport.json' reporter.
00:15:52 [0] [I] [Simulation] Finalized 'PropertyReport.json' reporter.
```

```

00:15:52 [0] [I] [Simulation] Finalizing 'SpatialReport' reporter.
00:15:52 [0] [I] [Simulation] Finalized 'SpatialReport' reporter.
00:15:52 [0] [I] [Simulation] Finalizing 'ReportEventRecorder.csv' reporter.
00:15:52 [0] [I] [Simulation] Finalized 'ReportEventRecorder.csv' reporter.
00:15:52 [0] [I] [Controller] Exiting execute_internal
00:15:52 [0] [I] [Eradication] Controller executed successfully.


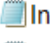
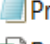
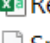

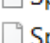
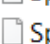
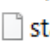
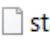
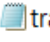

```

2.2.2 Cases: mpirun -n 2 & #SBATCH --ntasks=2

Results:

- Run succeeded.
- Less time.
- Different output.

Example output

			4/17/2024 12:19:40 PM
 InsetChart.json	7,866 KB		4/17/2024 12:29:36 PM
 PropertyReport.json	1,027 KB		4/17/2024 12:29:36 PM
 ReportEventRecorder.csv	9,209 KB		4/17/2024 12:29:35 PM
 SpatialReport_Adult_Vectors.bin	257 KB		4/17/2024 12:29:36 PM
 SpatialReport_Daily_Bites_Per_Human.bin	257 KB		4/17/2024 12:29:36 PM
 SpatialReport_Population.bin	257 KB		4/17/2024 12:29:36 PM
 SpatialReport_Prevalence.bin	257 KB		4/17/2024 12:29:36 PM
 state-21900-000.dtk	41,513 KB		4/17/2024 12:29:36 PM
 state-21900-001.dtk	41,478 KB		4/17/2024 12:29:36 PM
 transitions.json	1 KB		4/17/2024 12:19:41 PM

stdout.txt

```

...
00:09:54 [1] [I] [Simulation] Update(): Time: 21900.0 Rank: 1 StatPop: 1676 Infected: 328
00:09:54 [0] [I] [SerializedPopulation] Writing state to 'output/state-21900-000.dtk'
00:09:54 [1] [I] [SerializedPopulation] Writing state to 'output/state-21900-001.dtk'
00:09:55 [0] [I] [Simulation] Finalizing 'InsetChart.json' reporter.
00:09:55 [0] [I] [Simulation] Finalized 'InsetChart.json' reporter.
00:09:55 [0] [I] [Simulation] Finalizing 'PropertyReport.json' reporter.
00:09:55 [0] [I] [Simulation] Finalized 'PropertyReport.json' reporter.
00:09:55 [0] [I] [Simulation] Finalizing 'SpatialReport' reporter.
00:09:55 [1] [I] [Controller] Exiting execute_internal
00:09:55 [1] [I] [Eradication] Controller executed successfully.
00:09:55 [0] [I] [Simulation] Finalized 'SpatialReport' reporter.
00:09:55 [0] [I] [Simulation] Finalizing 'ReportEventRecorder.csv' reporter.
00:09:55 [0] [I] [Simulation] Finalized 'ReportEventRecorder.csv' reporter.

```


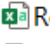
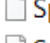
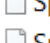
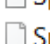
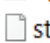
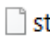
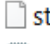
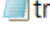

00:09:55 [0] [I] [Controller] Exiting execute_internal
 00:09:55 [0] [I] [Eradication] Controller executed successfully.

2.2.3 Cases: mpirun -n 3 & #SBATCH --ntasks=3

Results

- Run failed.
- Different output.

Example output

			4/17/2024 12:38:20 PM
	ReportEventRecorder.csv	9,081 KB	4/17/2024 12:47:59 PM
	SpatialReport_Adult_Vectors.bin	257 KB	4/17/2024 12:47:57 PM
	SpatialReport_Daily_Bites_Per_Human.bin	257 KB	4/17/2024 12:47:57 PM
	SpatialReport_Population.bin	257 KB	4/17/2024 12:47:57 PM
	SpatialReport_Prevalence.bin	257 KB	4/17/2024 12:47:57 PM
	state-21900-000.dtk	41,308 KB	4/17/2024 12:48:01 PM
	state-21900-001.dtk	40,659 KB	4/17/2024 12:48:01 PM
	state-21900-002.dtk	23,857 KB	4/17/2024 12:48:00 PM
	transitions.json	1 KB	4/17/2024 12:38:20 PM

stdout.txt

...
 00:09:41 [0] [E] [Eradication]

IllegalOperationException:
 Exception in baseReportLib/ChannelDataMap.cpp at 248 in Reduce.
 Channel=Variant Fraction-PfEMP1 Major from rank=2 had size=16331 while this rank=0 had size=21900

...
 00:09:41 [1] [E] [Eradication]

IllegalOperationException:
 Exception in baseReportLib/ChannelDataMap.cpp at 248 in Reduce.
 Channel=Variant Fraction-PfEMP1 Major from rank=2 had size=16331 while this rank=1 had size=21900

...
 00:09:41 [2] [E] [Eradication]

IllegalOperationException:

Exception in baseReportLib/ChannelDataMap.cpp at 248 in Reduce.

Channel=Variant Fraction-PfEMP1 Major from rank=0 had size=21900 while this rank=2 had size=16331

2.2.4 Cases: mpirun -n 4 & #SBATCH --ntasks=4

Results:

- Run failed.
- No output.

stdout.txt

...

00:00:01 [1] [I] [Simulation] Rank map contents not displayed until NodeRankMap::ToString() (re)implemented.

00:00:01 [2] [I] [Simulation] Merged rank 2 map now has 3 nodes.

00:00:01 [2] [I] [Simulation] Rank map contents not displayed until NodeRankMap::ToString() (re)implemented.

00:00:01 [2] [I] [Simulation] Initialized 'InsetChart.json' reporter

00:00:01 [3] [I] [Simulation] Merged rank 3 map now has 3 nodes.

00:00:01 [3] [I] [Simulation] Rank map contents not displayed until NodeRankMap::ToString() (re)implemented.

00:00:01 [3] [W] [Simulation] Rank 3 wasn't assigned any nodes! (# of procs is too big for simulation?)

00:00:01 [3] [I] [Eradication] Controller execution failed, exiting.

00:00:01 [0] [I] [Simulation] Merged rank 0 map now has 3 nodes.

00:00:01 [0] [I] [Simulation] Rank map contents not displayed until NodeRankMap::ToString() (re)implemented.

00:00:01 [0] [I] [Simulation] Initialized 'InsetChart.json' reporter

00:00:01 [0] [I] [Simulation] Initialized 'PropertyReport.json' reporter

00:00:01 [1] [I] [Simulation] Initialized 'InsetChart.json' reporter

00:00:01 [1] [I] [Simulation] Initialized 'PropertyReport.json' reporter

00:00:01 [2] [I] [Simulation] Initialized 'PropertyReport.json' reporter

00:00:01 [2] [I] [Simulation] 00:00:01 [0] [I] [Simulation] Initialized 'SpatialReport' reporter

00:00:01 [1] [I] [Simulation] Initialized 'SpatialReport' reporter

00:00:01 [1] [I] [Simulation] Initialized 'ReportEventRecorder.csv' reporter

00:00:01 [1] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...

00:00:01 [1] [I] [SimulationEventContext] 1 node(s) visited.

00:00:01 [1] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...

00:00:01 [1] [I] [SimulationEventContext] 1 node(s) visited.

00:00:01 [1] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...

00:00:01 [1] [I] [SimulationEventContext] 1 node(s) visited.

00:00:01 [1] [I] [StandardEventCoordinator] UpdateNodes() distributed

'NodeLevelHealthTriggeredIV' intervention to node 2


```

00:00:01 [1] [I] [StandardEventCoordinator] UpdateNodes() distributed
'NodeLevelHealthTriggeredIV' intervention to node 2
00:00:01 [1] [I] [StandardEventCoordinator] UpdateNodes() distributed
'NodeLevelHealthTriggeredIV' intervention to node 2
Initialized 'SpatialReport' reporter
00:00:01 [2] [I] [Simulation] Initialized 'ReportEventRecorder.csv' reporter
00:00:01 [2] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...
00:00:01 [2] [W] [SimulationEventContext] No nodes were visited. No nodes were added to the
event coordinator.
00:00:01 [2] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...
00:00:01 [2] [W] [SimulationEventContext] No nodes were visited. No nodes were added to the
event coordinator.
00:00:01 [2] [I] [SimulationEventContext] Time for campaign event. Calling Dispatch...
00:00:01 [2] [W] [SimulationEventContext] No nodes were visited. No nodes were added to the
event coordinator.

```

2.3 Special test

This is a normal test without invoking MPI manually, but with more SLURM directives setup:

```

#SBATCH --nodes=2
#SBATCH --ntasks=3
#SBATCH --cpus-per-task=2
#SBATCH --ntasks-per-core=3
#SBATCH --mem-per-cpu=8192

```

```

[zdf1921@quser34 ]$ sacct -X -P --format='jobid,state,elapsed,reqmem,reqcpus,ncpus' -j
8662087

```

```

JobID|State|Elapsed|ReqMem|ReqCPUS|NCPUS
8662087|COMPLETED|00:13:47|48G|6|6

```

output

..		4/19/2024 1:00:39 PM
InsetChart.json	7,862 KB	4/19/2024 1:14:21 PM
PropertyReport.json	1,027 KB	4/19/2024 1:14:21 PM
ReportEventRecorder.csv	9,079 KB	4/19/2024 1:14:18 PM
SpatialReport_Adult_Vectors.bin	257 KB	4/19/2024 1:14:21 PM
SpatialReport_Daily_Bites_Per_Human.bin	257 KB	4/19/2024 1:14:21 PM
SpatialReport_Population.bin	257 KB	4/19/2024 1:14:21 PM
SpatialReport_Prevalence.bin	257 KB	4/19/2024 1:14:21 PM
state-21900.dtk	58,360 KB	4/19/2024 1:14:20 PM
transitions.json	1 KB	4/19/2024 1:00:39 PM

Note: there is no big performance difference when running single simulation with just default settings.

3 Summary

Based on the testing results, we get a feeling that

- Seems like Eradication works differently with MPI cores #.
Q1: what do we want to achieve with MPI cores in EMOD?
Q2: should the results depend on (or be affected by) the MPI cores?
Q3: how to use the results if the results got affected by MPI cores #?
- Not sure User's suggestion is acceptable.
- We can't improve single simulation performance dramatically. Perhaps we may gain some overall performance for all simulations execution by choosing more nodes, mem, or cpus,...

What do you think?